

COL106 : 2021-22 (Semester I)

Project: Module 5

Pratik Kedia

Venkata Koppula

Akshay Mattoo

October 9, 2021

Notations

For two strings a, b , we denote $a.\text{concat}(b)$ by $a + b$ (that is, if $a = \text{"Hello"}$, and $b = \text{"World"}$, then $a + b = \text{"HelloWorld"}$). For any natural number n , $[n]$ denotes the set $\{1, 2, \dots, n\}$.

Instructions The lab-submission problem (marked ♠) is to be submitted via Moodle by 11:59PM on October 15th.

1 Introduction

The Central Board of Secondary Education (CBSE) has announced a new project – Academic BlockChain Document (ABCD) for providing transparent, tamper proof and paperless certificates (see [here](#) and other similar news reports). Inspired by this, our state **DS-Pradesh** would like to have a similar (maybe better?) solution for maintaining academic records – let’s call it ABCD++.

On a more serious note, the goal of this assignment is to illustrate the power of ideas developed in the previous lab-modules. Authenticated lists, authenticated sets play a crucial role in any cryptocurrency, but they have lots of other applications too, and this lab module will illustrate one such application. Additionally, we will also see some nice features of tree-based data structures.

Consider our hypothetical state **DS-Pradesh**. Every year, it has a number of students, whose scores need to be saved securely. Since we have already seen Merkle trees, we know that they provide a secure storage system, along with fast updates, and a mechanism to prove membership. We will use a combination of authenticated lists (module 2) and Merkle trees (module 3) to generate a secured storage system for all students’ scores. Additionally, if a student wishes to prove that he/she received a certain score in this exam, then he/she can provide a succinct proof, which can be verified efficiently.

In our system ABCD++, we will have an Authenticated List as the main structure, and each node in the list will store all students’ scores from a specific year. The students and their score data itself will be stored in the form of a Merkle tree, so you can assume that the number of students in any year is a power of 2.

1.1 Merkle Tree

The way the Merkle tree will be built is as follows - you will be given a list containing pairs of (student name, their score). We will assume for simplicity that the student id of each student is the index of its entry in this list. The students data will be stored in the leaf nodes of the tree, in the order of their student id. The `val` string of each leaf node will just contain the name of the student followed by an underscore and then his/her score (`<name>_<score>`). Apart from this, each `TreeNode` also has a `maxleafval` integer, which is maintained as the maximum score among all leaves of the subtree rooted at that `TreeNode`. For the leaf nodes, it is just the score of the student.

For all non-leaf nodes, the `val` string is the output of the CRF applied on the `val` strings of the two child nodes separated by a "#". For each node, you also need to maintain the `maxleafval` parameter and the `numberLeaves` parameter (it is the total number of leaves in the subtree rooted in that `TreeNode`).

1.2 Blockchain

The blockchain is an authenticated list (almost same as the one discussed in module 2). Each block in the blockchain contains pointers to the previous and next block, and strings `value` and `dgst`. In addition, every block also contains an integer year, representing the year of which the scores and student data is present. Also, each block has a Merkle tree, where as discussed above, the leaf nodes store all the students name and scores. The `value` of the block is calculated as follows - the Merkle tree corresponding to the block has a string `val` and an integer `maxleafval` associated with it. We concatenate the two attributes with an underscore in between, and get the `value` for that block. Thus `value = <mtree.rootnode.val>_<mtree.rootnode.maxleafval>`. Finally, the `dgst` for any block in the blockchain is computed by applying the CRF on the `dgst` of the previous node, and the `value` of that node (concatenated using #).

The DS-Pradesh State Board will only store the last block's `dgst` on its website.

1.3 Digital Certificates

Each student receives a *digital certificates* together with his/her final score. A digital certificate is simply a sibling-coupled-path-to-root in the Merkle tree. This digital certificate can be used by the student to prove that he/she got score x in the year y . Here is how the student will prove authenticity of his/her score:

- the student provides the digital certificate (corresponding to year y). The digital certificate ties his/her score x to the `val` of the Merkle tree root.
- next, the student provides a sequence of tuples $(\text{dgst}_y, \text{value}_y), (\text{dgst}_{y+1}, \text{value}_{y+1}), \dots, (\text{dgst}_{\text{curr}}, \text{value}_{\text{curr}})$, where dgst_j (resp. value_j) corresponds to the `dgst` (resp. `value`) for the block corresponding to year j . This sequence ties `val` to the last digest `dgstcurr`, and therefore it is not possible for someone to alter his/her scores (without breaking the CRF security).

Anyone can verify this proof by first checking the sibling-coupled-path-to-root, then checking that `val` is a substring of `valuey` (that is, the part before the "_"), then checking that each `dgstj` is correctly computed using `dgstj-1` and `valuej`, and finally, checking that `dgstcurr` is the latest `dgst` (available on the DS-Pradesh State Board website).

One can even prove that he/she scored the maximum marks in a particular year. We can also augment the Merkle tree with additional attributes so that a student can prove that he/she was in the top 5%.

1.4 Updating Scores

The scores in the latest block can get updated (due to regrade requests). Suppose the i^{th} student's score is updated. The update function first updates the `val` and `maxleafval` entries in the i^{th} leaf node. Next, it updates all the ancestors of this leaf node. Finally, it returns the `val` of the Merkle tree root.

2 Assignment Questions

You are given the following classes:

- `public class Pair<A,B>` : Objects of this class would be used to store each element in the sequence *Sibling-Coupled Path to Root*. The attributes of the class are:
 - `public A First`: the first element stored in the object.

- `public B Second`: the second element stored in the object.
- `public class TreeNode`: This class has the following attributes:
 - `public TreeNode parent`: the parent node
 - `public TreeNode left`: the left child node
 - `public TreeNode right`: the right child node
 - `public String val`: the value contained in this node
 - `public boolean isLeaf`: indicates whether the node is a leaf node or internal node
 - `public int numberLeaves`: the number of leaves in the sub-tree corresponding to this `TreeNode`.
 - `public int maxleafval`: the maximum score among all leaves of the subtree rooted at this `TreeNode`.
- `public class MerkleTree`: This class has the following attributes:
 - `public TreeNode rootnode`: pointer to the root node of the Merkle Tree
 - `public int numstudents`: number of students (n) whose records are stored in this Merkle tree.

This class has the following methods (to be implemented):

- `public String Build (List<Pair<String,int>> documents)`: Accepts an array of n documents, each of which is an `(String, int)` pair containing the name of a student (which is unique) and a non-negative integer corresponding to his/her score. It then builds a Merkle Tree using it, where the leaf node stores the string `<name>_<score>`. Eg: If the name was "Pete" and score was 42, string in `val` will be "Pete.42". Also, the `maxleafval` for that leaf node will be 42. For all other non-leaf nodes, `val` needs to be calculated as before, and `maxleafval` should be maintained as the maximum score among all leaves of the subtree rooted at that `TreeNode`. It returns the `val` corresponding to the root node.
The method also sets `numstudents = n`. You can assume n is a power of 2.
- `public String UpdateDocument(int student_id, int newScore)`: Updating a student's score given the student id¹ and newScore will be similar as before. Recall that only the scores in the final block are allowed to be updated.
First, go to the corresponding leaf node using the `student_id` and then change the `val` and `maxleafval` according to the new score. For all the ancestors of that leaf, the `val` and `maxleafval` need to change accordingly. Finally, return `val` of the `rootnode`.
- `public class Block`: This class has the following attributes:
 - `public Block previous`: pointer to previous block.
 - `public Block next`: pointer to next block.
 - `public String dgst`: a string representing the digest of this block.
 - `public MerkleTree mtree`: the Merkle Tree contained in the current Block of the `BlockChain`.
 - `public int year`: the year for which data is stored in this block.
 - `public String value`: it is the string defined as follows:
`value= mtree.rootnode.val+ "_" + mtree.rootnode.maxleafval`

¹Recall, the student id is the position of this student in the Merkle tree leaves

- **public class Blockchain:** This class has the following attributes:
 - **public static final String start_string:** a string representing the starting digest for all objects of this class.
This string should be equal to “LabModule5”.
 - **public Block lastblock:** represents the last block present in the blockchain.
 - **public Block firstblock:** represents the first block present in the blockchain.

The class has the following methods (to be implemented):

- **public String InsertBlock(List<Pair<String,int>> Documents, int inputyear):** first it creates a new **Block** with the year corresponding to inputyear and Merkle tree built using documents. It then adds this new block to the blockchain. Thereafter, it computes the **dgst** using the **value** and the previous **dgst** as follows:

$$u.dgst = \begin{cases} CRF64.Fn(BlockChain.start_string + \# + u.value) & \text{if } u.previous = null \\ CRF64.Fn(u.previous.dgst + \# + u.value) & \text{if } u.previous \neq null \end{cases}$$

where *CRF64* is an instance of the class *CRF* with **outputsize = 64**.

Finally, it returns the digest of the last **Block** of the updated **Blockchain**.

- **public Pair<List<Pair<String,String>>,List<Pair<String,String>>> ProofofScore (int student_id, int year):** Accepts the id of student and returns a pair. The first element in the pair is *Sibling-Coupled Path to Root* of the corresponding Merkle tree node. The second element in the pair is the list of pair of values and digest (n.value, n.dgst) of all nodes of the **Blockchain**, from the block corresponding to input year to the last block (both blocks included).

2.1 Exercises

1. Merkle Tree - Proof of membership and updates:

Implementing methods of class **MerkleTree**

- ♠ **Build:** This method accepts an array of n students data and builds a Merkle Tree using it.
 - the index/id of a student in a particular year is the index of the list where the input data is received.²
 - the input contains pair of string and integers, the string is the student’s name and the corresponding integer is his/her score for that year.
 - Each of these students data would be stored at the leaf nodes.
 - Each of the non leaf nodes would store the CRF output after concatenating the strings stored of the two child nodes.
 - At the end, the root of the tree should be stored at **rootnode** and n (length of **documents**) should be stored at **numstudents**.
- ♠ **UpdateDocument:** Takes input $student_id \in [0, n - 1]$ and **newScore** and updates the corresponding student record.
 - Start from $N_{0,1}$ and traverse down the Merkle Tree until you reach $N_{\log(n), student_id}$
 - Update the score part in the string $val_{\log(n), student_id}$ with **newScore**. Also update the **maxleafval** of that node.
 - Traverse back up from $N_{\log(n), student_id}$ to the root $N_{0,1}$ while updating $val_{i,j}$ and **maxleafval** (if needed) for each ancestor node $N_{i,j}$ encountered in the path.

²Note the array indices of the input array are from 0 to $n - 1$.

- Return `val` of the `rootnode` as the updated `summary`.

2. BlockChain:

Implementing methods of `BlockChain` :

- InsertBlock**: this method takes as input a list of student records and an integer `inputyear`.
 - It first creates a new object of class `Block`. It sets the pointer to previous and next blocks appropriately and sets the attribute `year` equal to the `inputyear` integer.
 - Then it builds a Merkle tree with the provided list and sets the `mtree` attribute to it. The Merkle tree `rootnode` is used to calculate the `value` of the `Block` as explained earlier. The `value` is then used to compute the `dgst` of the newly inserted `Block` using the `dgst` of the immediately preceding `Block` (or `start_string`, as explained earlier).
 - Finally this `Block` object is appended to the `BlockChain` and the `dgst` of the last `Block` of the updated `BlockChain` is returned.
- ProofofScore**: It takes as input a `student_id` and an `year`. It then returns a pair to retrieve and verify the student's score.
 - It iterates over the list to find the block corresponding to the input year.
 - It goes to the corresponding `mtree` node using the `student_id`.
 - It collects the *Sibling-Coupled Path to Root* of the corresponding Merkle tree node in the form of a list of pairs.
 - It collects the values and digests of all nodes, starting from the block corresponding to year, to the last block, in the form of a list of pairs.
 - Finally, it returns a pair of the two lists collected. This list can be used to know and authenticate the score of a student in a particular year.

3 Instructions

- Do not change the accessibility, names or signatures of the attributes and methods in the driver code. You may add your own attributes and methods to the `Block`, `BlockChain` and `MerkleTree` classes.
- The default constructor is used to instantiate objects of all the above classes. It is your responsibility to ensure appropriate initialization of the attributes of a newly created object.
- **Submission instructions for lab-submission problem:** You must submit `BlockChain.java`, `Block.java` and `MerkleTree.java` on Moodle. You must create a directory whose name is your Kerberos id, followed by "Module5" (for example, if your Kerberos id is "xyz120100", then the folder name should be "xyz120100Module5"). The directory must contain `MerkleTree.java`, `Block.java` and `BlockChain.java`. Finally, compress this directory, and upload it on Moodle. The file name should be `kerberosidModule5.zip` (that is, `xyz120100Module5.zip` in the above example).

Acknowledgements

This is Version 1 of the document. Many thanks to the following students for identifying typos in previous version: ³ Simran Malik, Tanish Singh Tak, Tushita Pandey, Vaibhav Misra, Somaditya Singh, Divyansh Mittal, Himasindhu Appili, Manas Jain, Prateek Mishra, Chirag, Aditya Mathur, Shivam Jain, Atif Anwar, Yash Pravin Shirke, Hemali Priyadarshi, Sachit Sachdeva, Dhruv Tyagi, Maitree Shandilya, Puurshottam Malviya, Yeruva Hitesh Reddy, Kushagra Rode, Sreemanti Dey, Adit Malhotra, Tanish Gupta, Abhinav Barnawal, Rishabh Dhiman, Vansh Kacchwal, Aditya Agrawal. If you find any errors, please send an email to kvenkata@iitd.ac.in and Akshay.Mattoo.me218@mech.iitd.ac.in (and participate in the 'error-finding competition').

³If you found an error and your name is not listed here, please send an email to kvenkata@iitd.ac.in.